

Bank customer

October 9, 2024

```
[35]: import pandas as pd
```

```
[37]: #Step 1: Read the Dataset - Load the dataset
```

```
[51]: df = pd.read_csv('/home/pc13/Downloads/archive/Churn_Modelling.csv')
```

```
[50]: df.shape
```

```
[50]: (10000, 14)
```

```
[52]: # Display the first few rows of the dataset
print(df.head())
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
[ ]: #Step 2: Distinguish Feature and Target Set
```

```
[39]: from sklearn.model_selection import train_test_split

# Define features and target
X = df.drop(columns=['Exited', 'CustomerId']) # 'Exited' is the target variable
y = df['Exited']

# Convert categorical variables to dummy/indicator variables
X = pd.get_dummies(X, drop_first=True)

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Output the shape of the datasets
print(f"Training set shape: {X_train.shape}, Test set shape: {X_test.shape}")
```

Training set shape: (8000, 2943), Test set shape: (2000, 2943)

```
[40]: #Step 3: Normalize the Train and Test Data
```

```
[41]: from sklearn.preprocessing import StandardScaler

# Initialize the StandardScaler
scaler = StandardScaler()

# Fit and transform the training data, transform the test data
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Output the mean and variance of the scaled training data
print(f"Mean of X_train: {X_train.mean(axis=0)}, Variance of X_train: {X_train.
↳var(axis=0)}")
```

Mean of X_train: [-1.59872116e-17 5.43565193e-16 -1.89626093e-16 ...
-4.26325641e-17
-7.19424520e-17 2.84217094e-17], Variance of X_train: [1. 1. 1. ... 1. 1. 1.]

```
[42]: #Step 4: Initialize and Build the Model
```

```
[43]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

# Build the model
model = keras.Sequential([
    layers.Dense(32, activation='relu', input_shape=(X_train.shape[1],)),
    layers.Dense(16, activation='relu'),
```

```

        layers.Dense(1, activation='sigmoid') # For binary classification
    ])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])

# Train the model
history = model.fit(X_train, y_train, epochs=50, batch_size=32,
    ↪validation_split=0.2)

# Output training history
print(history.history)

```

```

/home/pc13/miniconda3/lib/python3.12/site-
packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
2024-10-09 13:36:57.898555: W
external/local_tsl/tsl/framework/cpu_allocator_impl.cc:83] Allocation of
75340800 exceeds 10% of free system memory.

```

```

Epoch 1/50
200/200          1s 3ms/step -
accuracy: 0.7058 - loss: 0.6054 - val_accuracy: 0.7944 - val_loss: 0.5141
Epoch 2/50
200/200          0s 2ms/step -
accuracy: 0.8088 - loss: 0.4159 - val_accuracy: 0.7713 - val_loss: 0.5427
Epoch 3/50
200/200          0s 2ms/step -
accuracy: 0.8347 - loss: 0.3509 - val_accuracy: 0.7531 - val_loss: 0.5723
Epoch 4/50
200/200          0s 2ms/step -
accuracy: 0.8639 - loss: 0.3021 - val_accuracy: 0.7369 - val_loss: 0.6044
Epoch 5/50
200/200          0s 2ms/step -
accuracy: 0.8704 - loss: 0.2827 - val_accuracy: 0.7287 - val_loss: 0.6332
Epoch 6/50
200/200          0s 2ms/step -
accuracy: 0.8831 - loss: 0.2568 - val_accuracy: 0.7181 - val_loss: 0.6699
Epoch 7/50
200/200          0s 2ms/step -
accuracy: 0.8852 - loss: 0.2541 - val_accuracy: 0.7175 - val_loss: 0.7098
Epoch 8/50
200/200          0s 2ms/step -
accuracy: 0.8959 - loss: 0.2343 - val_accuracy: 0.7244 - val_loss: 0.7259
Epoch 9/50

```

200/200 0s 2ms/step -
 accuracy: 0.8968 - loss: 0.2255 - val_accuracy: 0.7119 - val_loss: 0.7891
 Epoch 10/50
 200/200 0s 2ms/step -
 accuracy: 0.9078 - loss: 0.2051 - val_accuracy: 0.7169 - val_loss: 0.7816
 Epoch 11/50
 200/200 1s 3ms/step -
 accuracy: 0.9125 - loss: 0.1910 - val_accuracy: 0.7119 - val_loss: 0.8578
 Epoch 12/50
 200/200 1s 2ms/step -
 accuracy: 0.9224 - loss: 0.1759 - val_accuracy: 0.7144 - val_loss: 0.8983
 Epoch 13/50
 200/200 0s 2ms/step -
 accuracy: 0.9317 - loss: 0.1634 - val_accuracy: 0.7206 - val_loss: 0.9857
 Epoch 14/50
 200/200 0s 2ms/step -
 accuracy: 0.9431 - loss: 0.1363 - val_accuracy: 0.7169 - val_loss: 1.0089
 Epoch 15/50
 200/200 0s 2ms/step -
 accuracy: 0.9505 - loss: 0.1248 - val_accuracy: 0.7262 - val_loss: 1.0446
 Epoch 16/50
 200/200 0s 2ms/step -
 accuracy: 0.9572 - loss: 0.1087 - val_accuracy: 0.7337 - val_loss: 1.1106
 Epoch 17/50
 200/200 0s 2ms/step -
 accuracy: 0.9613 - loss: 0.0965 - val_accuracy: 0.7269 - val_loss: 1.1864
 Epoch 18/50
 200/200 0s 2ms/step -
 accuracy: 0.9681 - loss: 0.0858 - val_accuracy: 0.7306 - val_loss: 1.2437
 Epoch 19/50
 200/200 0s 2ms/step -
 accuracy: 0.9740 - loss: 0.0753 - val_accuracy: 0.7275 - val_loss: 1.3100
 Epoch 20/50
 200/200 0s 2ms/step -
 accuracy: 0.9775 - loss: 0.0678 - val_accuracy: 0.7312 - val_loss: 1.3680
 Epoch 21/50
 200/200 1s 3ms/step -
 accuracy: 0.9803 - loss: 0.0566 - val_accuracy: 0.7344 - val_loss: 1.4062
 Epoch 22/50
 200/200 0s 2ms/step -
 accuracy: 0.9859 - loss: 0.0445 - val_accuracy: 0.7375 - val_loss: 1.4660
 Epoch 23/50
 200/200 0s 2ms/step -
 accuracy: 0.9879 - loss: 0.0416 - val_accuracy: 0.7331 - val_loss: 1.5544
 Epoch 24/50
 200/200 0s 2ms/step -
 accuracy: 0.9897 - loss: 0.0339 - val_accuracy: 0.7369 - val_loss: 1.6068
 Epoch 25/50

200/200 0s 2ms/step -
 accuracy: 0.9919 - loss: 0.0298 - val_accuracy: 0.7337 - val_loss: 1.6703
 Epoch 26/50
 200/200 0s 2ms/step -
 accuracy: 0.9929 - loss: 0.0315 - val_accuracy: 0.7375 - val_loss: 1.7006
 Epoch 27/50
 200/200 0s 2ms/step -
 accuracy: 0.9940 - loss: 0.0227 - val_accuracy: 0.7369 - val_loss: 1.7706
 Epoch 28/50
 200/200 0s 2ms/step -
 accuracy: 0.9931 - loss: 0.0262 - val_accuracy: 0.7412 - val_loss: 1.8225
 Epoch 29/50
 200/200 0s 2ms/step -
 accuracy: 0.9967 - loss: 0.0178 - val_accuracy: 0.7394 - val_loss: 1.8540
 Epoch 30/50
 200/200 0s 2ms/step -
 accuracy: 0.9967 - loss: 0.0179 - val_accuracy: 0.7437 - val_loss: 1.9330
 Epoch 31/50
 200/200 0s 2ms/step -
 accuracy: 0.9971 - loss: 0.0161 - val_accuracy: 0.7400 - val_loss: 1.9963
 Epoch 32/50
 200/200 1s 3ms/step -
 accuracy: 0.9964 - loss: 0.0142 - val_accuracy: 0.7425 - val_loss: 2.0427
 Epoch 33/50
 200/200 1s 3ms/step -
 accuracy: 0.9967 - loss: 0.0127 - val_accuracy: 0.7381 - val_loss: 2.0862
 Epoch 34/50
 200/200 1s 3ms/step -
 accuracy: 0.9969 - loss: 0.0145 - val_accuracy: 0.7494 - val_loss: 2.1251
 Epoch 35/50
 200/200 1s 3ms/step -
 accuracy: 0.9973 - loss: 0.0106 - val_accuracy: 0.7519 - val_loss: 2.1650
 Epoch 36/50
 200/200 0s 2ms/step -
 accuracy: 0.9967 - loss: 0.0112 - val_accuracy: 0.7519 - val_loss: 2.2162
 Epoch 37/50
 200/200 0s 2ms/step -
 accuracy: 0.9979 - loss: 0.0116 - val_accuracy: 0.7431 - val_loss: 2.2667
 Epoch 38/50
 200/200 0s 2ms/step -
 accuracy: 0.9981 - loss: 0.0087 - val_accuracy: 0.7481 - val_loss: 2.3361
 Epoch 39/50
 200/200 0s 2ms/step -
 accuracy: 0.9983 - loss: 0.0081 - val_accuracy: 0.7487 - val_loss: 2.3949
 Epoch 40/50
 200/200 0s 2ms/step -
 accuracy: 0.9964 - loss: 0.0119 - val_accuracy: 0.7513 - val_loss: 2.4292
 Epoch 41/50

200/200 0s 2ms/step -
accuracy: 0.9970 - loss: 0.0086 - val_accuracy: 0.7456 - val_loss: 2.4132
Epoch 42/50

200/200 1s 3ms/step -
accuracy: 0.9979 - loss: 0.0063 - val_accuracy: 0.7456 - val_loss: 2.4626
Epoch 43/50

200/200 1s 2ms/step -
accuracy: 0.9984 - loss: 0.0065 - val_accuracy: 0.7494 - val_loss: 2.4742
Epoch 44/50

200/200 0s 2ms/step -
accuracy: 0.9990 - loss: 0.0053 - val_accuracy: 0.7519 - val_loss: 2.6020
Epoch 45/50

200/200 0s 2ms/step -
accuracy: 0.9992 - loss: 0.0054 - val_accuracy: 0.7506 - val_loss: 2.5558
Epoch 46/50

200/200 0s 2ms/step -
accuracy: 0.9980 - loss: 0.0092 - val_accuracy: 0.7419 - val_loss: 2.6279
Epoch 47/50

200/200 0s 2ms/step -
accuracy: 0.9971 - loss: 0.0097 - val_accuracy: 0.7594 - val_loss: 2.5883
Epoch 48/50

200/200 0s 2ms/step -
accuracy: 0.9959 - loss: 0.0160 - val_accuracy: 0.7525 - val_loss: 2.5387
Epoch 49/50

200/200 1s 3ms/step -
accuracy: 0.9944 - loss: 0.0191 - val_accuracy: 0.7456 - val_loss: 2.5016
Epoch 50/50

200/200 0s 2ms/step -
accuracy: 0.9972 - loss: 0.0116 - val_accuracy: 0.7563 - val_loss: 2.5447
{'accuracy': [0.7549999952316284, 0.8125, 0.8364062309265137,
0.8528125286102295, 0.8665624856948853, 0.8732812404632568, 0.8785937428474426,
0.8860937356948853, 0.8910937309265137, 0.9014062285423279, 0.9059374928474426,
0.91796875, 0.9262499809265137, 0.9315624833106995, 0.942187488079071,
0.9496874809265137, 0.9573437571525574, 0.9637500047683716, 0.9714062213897705,
0.9765625, 0.9801562428474426, 0.9829687476158142, 0.9870312213897705,
0.9871875047683716, 0.9898437261581421, 0.9917187690734863, 0.9921875,
0.9937499761581421, 0.9950000047683716, 0.9964062571525574, 0.9965624809265137,
0.9962499737739563, 0.9957812428474426, 0.99609375, 0.9971874952316284,
0.9965624809265137, 0.9976562261581421, 0.9982812404632568, 0.9981250166893005,
0.9971874952316284, 0.9981250166893005, 0.9976562261581421, 0.9984375238418579,
0.999218761920929, 0.9989062547683716, 0.9975000023841858, 0.9962499737739563,
0.9959375262260437, 0.9950000047683716, 0.996874988079071], 'loss':
[0.5601981282234192, 0.412230521440506, 0.34983593225479126,
0.31614336371421814, 0.29304665327072144, 0.27458441257476807,
0.2608238160610199, 0.24377702176570892, 0.23097918927669525,
0.21321630477905273, 0.1997666358947754, 0.18386124074459076,
0.1677486002445221, 0.15334337949752808, 0.13562412559986115,
0.11958098411560059, 0.104751355946064, 0.09151662141084671,

```

0.07862219214439392, 0.0673275962471962, 0.058343637734651566,
0.04943905770778656, 0.04263339564204216, 0.037815384566783905,
0.03324383497238159, 0.03129428252577782, 0.027008483186364174,
0.024421803653240204, 0.02057233266532421, 0.017779359593987465,
0.01656091958284378, 0.015143761411309242, 0.014990261755883694,
0.015423358418047428, 0.011869342066347599, 0.011964371427893639,
0.010659548453986645, 0.008586808107793331, 0.008896760642528534,
0.010534089989960194, 0.008305300958454609, 0.007050278130918741,
0.0066164471209049225, 0.004999118857085705, 0.005917804781347513,
0.009906486608088017, 0.013881074264645576, 0.016141420230269432,
0.016124999150633812, 0.012920962646603584], 'val_accuracy':
[0.7943750023841858, 0.7712500095367432, 0.753125011920929, 0.7368749976158142,
0.7287499904632568, 0.7181249856948853, 0.7174999713897705, 0.7243750095367432,
0.7118750214576721, 0.7168750166893005, 0.7118750214576721, 0.7143750190734863,
0.7206249833106995, 0.7168750166893005, 0.7262499928474426, 0.7337499856948853,
0.7268750071525574, 0.7306249737739563, 0.7275000214576721, 0.731249988079071,
0.734375, 0.737500011920929, 0.7331249713897705, 0.7368749976158142,
0.7337499856948853, 0.737500011920929, 0.7368749976158142, 0.7412499785423279,
0.7393749952316284, 0.7437499761581421, 0.7400000095367432, 0.7425000071525574,
0.7381250262260437, 0.7493749856948853, 0.7518749833106995, 0.7518749833106995,
0.7431250214576721, 0.7481250166893005, 0.7487499713897705, 0.7512500286102295,
0.7456250190734863, 0.7456250190734863, 0.7493749856948853, 0.7518749833106995,
0.7506250143051147, 0.7418749928474426, 0.7593749761581421, 0.7524999976158142,
0.7456250190734863, 0.7562500238418579], 'val_loss': [0.5140730142593384,
0.5426977276802063, 0.5722506046295166, 0.6044119596481323, 0.6331974864006042,
0.6698675751686096, 0.709787905216217, 0.7259223461151123, 0.7891015410423279,
0.7815655469894409, 0.8578264713287354, 0.8983173966407776, 0.9856798052787781,
1.0088597536087036, 1.044642448425293, 1.1106300354003906, 1.1864399909973145,
1.243656039237976, 1.3099699020385742, 1.3680219650268555, 1.406174898147583,
1.466016411781311, 1.55442476272583, 1.6068129539489746, 1.6702532768249512,
1.7006173133850098, 1.7706114053726196, 1.8225336074829102, 1.8539929389953613,
1.933018445968628, 1.9963420629501343, 2.042693853378296, 2.086225748062134,
2.1250882148742676, 2.1650307178497314, 2.2161989212036133, 2.266688823699951,
2.3361470699310303, 2.3949286937713623, 2.4292209148406982, 2.4132468700408936,
2.4626195430755615, 2.4742307662963867, 2.6020114421844482, 2.5558199882507324,
2.627865791320801, 2.5882883071899414, 2.5387377738952637, 2.501645803451538,
2.5447041988372803]}

```

[47]: *# Neural Network Training Progress Overview*

Key Components

- ****Epochs****: Each epoch indicates one complete cycle through the training dataset. In this output, the model was trained **for** 50 epochs.
- ****Batch Progress****: The `200/200` indicates that the training process has completed 200 batches **for** that epoch.

```

- Metrics:
  - accuracy: The accuracy of the model on the training set for that epoch.
  - loss: The training loss, which measures how well the model is
    performing (lower is better).
  - val_accuracy: The accuracy of the model on the validation set (data not
    seen by the model during training).
  - val_loss: The loss on the validation set.

## Observations

- Training Accuracy: Starts around 70.58% in the first epoch and improves
  to about 99.22% by the 50th epoch, indicating the model is learning well.

- Training Loss: Starts at 0.6054 and decreases to 0.0054, showing that the
  model is reducing its error on the training data.

- Validation Metrics:
  - Validation accuracy starts at 79.44% and fluctuates, eventually reaching 75.
    63%. The slight drop or fluctuation in validation accuracy compared to
    training accuracy can indicate overfitting, where the model learns the
    training data too well but does not generalize effectively to unseen data.
  - Validation loss increases initially before decreasing but stabilizes,
    suggesting the model's performance on unseen data may not improve beyond a
    certain point.

```

Cell In[47], line 17

```

- Training Accuracy: Starts around 70.58% in the first epoch and
  improves to about 99.22% by the 50th epoch, indicating the model is learning
  well.

^
SyntaxError: invalid decimal literal

```

[]: #Step 5: Evaluate the Model

```

[33]: from sklearn.metrics import accuracy_score, confusion_matrix

# Make predictions
y_pred = (model.predict(X_test) > 0.5).astype("int32")

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

```



```
# Print confusion matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)
```

```
63/63          0s 2ms/step
Accuracy: 0.74
Confusion Matrix:
[[1360  247]
 [ 270  123]]
```

```
[ ]:
```

```
[ ]:
```